

QUESTION CLASSIFIER

MATEUSZ DOROBK



The background of the slide is a light gray gradient. It is decorated with numerous realistic water droplets of various sizes. Some droplets are at the top left, some are scattered in the middle, and a larger cluster is at the bottom right. Each droplet has a highlight and a shadow, giving it a three-dimensional appearance.

THE PROBLEM

Data comes in `.xml` format, but better to read it from csv file.

```
df = pd.read_csv("Questions.csv", sep=";")  
df.head()
```

Question	Type
Is Hirschsprung disease a mendelian or a multifactorial disorder?	summary
List signaling molecules (ligands) that interact with the receptor EGFR?	list
Is the protein <u>Papilin</u> secreted?	<u>yesno</u>
Are long non coding RNAs spliced?	<u>yesno</u>
Is RANKL secreted from the cells?	<u>yesno</u>

DATA

For this problem we should use `multiclass` classifier.

```
print(df.Type.unique())  
>> ['summary' 'list' 'yesno' 'factoid']
```

```
df.shape  
>> (2251, 2)
```

DATA


PREPROCESSING



First step of preprocessing string is to tokenize it - change each word into separate string and gather them into a list. I've used `nltk` method which has some additional features for example separates punctuation to different tokens.

```
tokens = nltk.word_tokenize("Is Hirschsprung disease a multifactorial  
disorder?")  
>> ['Is', 'Hirschsprung', 'disease', 'a', 'multifactorial', 'disorder', '?']
```

TOKENIZING



In every language there are stop words, that actually don't give much information in a sentence. For example (a, the, in, or, ...).

```
tokens = [token for token in tokens if token not in stopwords_en]  
>> ['Is', 'Hirschsprung', 'disease', 'mendelian', 'multifactorial',  
     'disorder', '?']
```

STOPWORDS



Removing punctuation also helps reduce number of tokens that not necessary increase informative value of sentence.

```
tokens = [token for token in tokens if token not in punctuation]
>> ['Is', 'Hirschsprung', 'disease', 'mendelian', 'multifactorial',  
    'disorder']
```

```
# ['123a45n6', 'example!', 'witho0ut', 'non-letters']
tokens = [re.sub(r'^a-zA-Z', '', token) for token in tokens]
>> ['an', 'example', 'without', 'nonletters']
```

PUNCTUATION

"Lemmatization (or lemmatization) in linguistics is the process of grouping together the inflected forms of a word so they can be analyzed as a single item, identified by the word's lemma, or dictionary form." (Wiki)

```
# ['list', 'signaling', 'molecules', 'ligands', 'interact', 'receptor']  
tokens = [lemmatize(pair) for pair in pos_tag(tokens)]  
>> ['list', 'signal', 'molecule', 'ligands', 'interact', 'receptor']
```

LEMMATIZATION

Stemming has basically the same purpose as Lemmatization, but is performed with regex rules, which makes it way faster, and sometimes allow to decrease number of unique tokens in dataset even after lemmatization.

```
# ['list', 'signal', 'molecule', 'ligands', 'interact', 'receptor']  
tokens = [porter.stem(token) for token in tokens]  
>> ['list', 'signal', 'molecul', 'ligand', 'interact', 'receptor']
```

STEMMING

I've used `sklearn.feature_extraction.text.CountVectorizer` to Vectorize my data. It takes text file as input but there is a short trick with `StringIO` that allows me to transform data to proper format.

```
with StringIO('\n'.join([i for i in questions.values])) as text:  
    count_vect = CountVectorizer(analyzer=preprocess_text)  
    count_vect.fit_transform(text)
```

In our dataset after preprocessing there are 3601 tokens (more than training examples) we will have to deal with it later.

```
len(count_vect.vocabulary_)  
>> 3601
```

VECTORIZATION

There is a vocabulary of words. As we can see in first example not all of them are regular words in english.

```
words_sorted_by_index, _ = zip(*sorted(count_vect.vocabulary_.items(),
key=itemgetter(1)))
words_sorted_by_index[:5]
>> ('aa', 'aagena', 'abacavir', 'abatacept', 'abc')
```

This is our final dataset shape, time to do the classification.

```
count_vect.transform([i for i in questions.values]).toarray().shape
>> (2251, 3601)
```

FINAL DATA SHAPE

CLASSIFICATION



Decision Tree

Accuracy Train: 99.8% Accuracy Valid: 70.4%

Best params: `class_weight = 'balanced'`, `presort = False`

	0	1	2	3
0	134	8	3	7
1	16	85	3	33
2	0	3	188	1
3	62	43	8	82

DECISION TREE

Random Forest Accuracy Train: 93.5% Accuracy Valid: 72.3%

Best params: `class_weight = 'balanced'`, `max_depth = 30`

`max_features = 19`, `n_estimators = 100`

	0	1	2	3
0	114	10	3	25
1	16	66	0	55
2	0	1	190	1
3	52	31	6	106

RANDOM FOREST

K-Neares Neighbours Accuracy Train: 99.8% Accuracy Valid: 32.1%
Best params: n_neighbors = 10, weights = 'distance'

	0	1	2	3
0	13	16	2	121
1	11	23	1	102
2	27	35	19	111
3	9	21	3	162

K-NEARES NEIGHBOURS

Logistic Regression

Accuracy Train: 85.7%

Accuracy Valid: 74.6%


Best params: C = 0.1, multi_class = 'multinomial', solver = 'lbfgs'

	0	1	2	3
0	134	8	3	7
1	14	67	1	55
2	1	0	189	2
3	55	19	7	114

LOGISTIC REGRESSION

NOTES

- I'VE TESTED THAT **PCA** DOESN'T IMPROVE PERFORMANCE OF ANY OF CLASSIFIERS.
- USING **STANDARSCALER()** WASN'T A GOOD DUE TO BINARY CHARACTER OF DATA.
- MY VALIDATION METRIC WAS **ACCURACY** DUE TO EVEN DISTIBUTION IN CLASS.



I've written short function to classify inputed by user questions to one of 4 classes.

```
def predict_question(question):  
    x = count_vect.transform([question]).toarray()  
    return classes[clf.predict(x)[0]]
```

RESULTS






RESULTS

- *LIST TWO OF YOUR FAVOURIE FILMS. LIST*
- *WHERE ARE YOU? FACTOID*
- *HOW OLD ARE YOU? SUMMARY*



RESULTS

- *DO YOU LIKE TO STUDY? YESNO*
 - *HOW DO YOU FEEL RRIGHT NOW? SUMMARY*
 - *WHAT IS YOUR NAME? SUMMARY*
- 

The background is a light gray gradient. It is decorated with several realistic water droplets of various sizes. Some droplets are at the top left, some are at the bottom right, and others are scattered in the center. Each droplet has a highlight and a shadow, giving it a three-dimensional appearance.

THANK YOU